

Python Libraries for Machine Learning

Useful libraries for us

- [numpy](#): data structures (esp. array and matrix) management
- [scipy](#): algorithms, scientific computing
- [matplotlib](#): graphs
- [pandas](#): data frame management
- [scikit-learn](#): machine learning, text processing
- [nltk](#): natural language processing

numpy

- N-dimensional array objects
- Linear algebra, Fourier transform, random number generation

First steps...

```
import numpy as np
```

```
A = np.array([[1, 2, 3], [4, 5, 6]])
```

```
print(A)
```

```
[[1 2 3]
```

```
 [4 5 6]]
```

```
Af = np.array([1, 2, 3], float)
```

Ranges

```
np.arange(0, 1, 0.2)
```

```
array([0. , 0.2, 0.4, 0.6, 0.8])
```

Linear model

```
#creates a linear model consisting of 4  
datapoints in the range  $[0, 2*\pi]$ 
```

```
np.linspace(0, 2*np.pi, 4)
```

```
array([0., 2.0943951 , 4.1887902 ,  
6.28318531])
```

Initializing matrices

```
A = np.zeros((2,3))
```

A

```
array([[0., 0., 0.],  
       [0., 0., 0.]])
```

- Try also `np.ones`, `np.diag`

Random

```
np.random.random( ( 2 , 3 ) )
```

```
array([[0.70837754, 0.66890078, 0.19124653],  
       [0.64323459, 0.85562671, 0.83175077]])
```


Saving/loading to/from file

```
np.savetxt("a_out.txt", a)
```

```
b = np.loadtxt("a_out.txt")
```

File format

0.000000000000000000e+00 0.000000000000000000e+00 0.000000000000000000e+00

0.000000000000000000e+00 0.000000000000000000e+00 0.000000000000000000e+00

Modifying arrays

```
A = np.zeros((2, 2))
```

```
A[0, 0] = 1
```

```
A
```

```
array([[1., 0.],  
       [0., 0.]])
```

Reshaping arrays

```
a = np.arange(10).reshape((2,5))
```

First it creates an array of 10 items, then it reshapes in a 2x5 matrix

Attributes

<code>a.ndim</code>	<i># returns the number of array dimensions (e.g., 2)</i>
<code>a.shape</code>	<i># returns shape of array (e.g., 2x5)</i>
<code>a.size</code>	<i># returns the number of of elements (e.g., 10)</i>
<code>a.T</code>	<i># transposes the array</i>
<code>a.dtype</code>	<i># returns the array data type</i>

Array operations

```
a = np.arange(4)  
# array([0, 1, 2, 3])
```

```
b = np.array([2, 3, 2, 4])
```

```
a*b  
#array([0, 3, 4, 12])
```

```
b - a  
# array([2, 2, 0, 1])
```

```
c = [2, 3, 4, 5]  
a*c  
#array([0, 3, 8, 15])
```

Operations with scalars

```
A = np.ones((3,3))
```

```
print(3 * A - 1)
```

```
[[2. 2. 2.]
```

```
 [2. 2. 2.]
```

```
 [2. 2. 2.]]
```

Matrix operations

- inner product
- outer product
- dot product (matrix multiplication)
- matrix element multiplication (matmul or * operator)
- **numpy automatically converts lists into numpy arrays**

Matrix operations

note: numpy automatically converts lists

`u = [1, 2, 3]`

`v = [1, 1, 1]`

`np.dot(u, v)`

#6

Matrix operations

```
A = np.ones((3, 2))
```

```
B = np.ones((2, 3))
```

```
np.dot(A, B)
```

```
array([[2., 2., 2.],  
       [2., 2., 2.],  
       [2., 2., 2.]])
```

Matrix operations

```
np.dot(B, A)
```

```
array([[3., 3.],  
       [3., 3.]])
```

Slicing arrays

`a[2, :]` *# third row, all columns*

`a[1:3]` *# 2nd, 3rd row, all columns*

`a[:, 2:4]` *# all rows, columns 3 and 4*

Iterating

- Over rows

```
for row in A:  
    print(row)
```

As flat values

```
for element in A.flat:  
    print(element)
```

1.0

1.0

1.0

1.0

1.0

1.0

Linear algebra

- `linalg.inv(A)` inverse matrix
- `linalg.solve(A,b)` solves the system $Ax=b$
- `linalg.svd(A, full)` singular value decomposition

scipy

Algorithms and mathematical tools built to work with NumPy arrays.

- linear algebra - [scipy.linalg](#)
- statistics - [scipy.stats](#)
- optimization - [scipy.optimize](#)
- sparse matrices - [scipy.sparse](#)
- signal processing - [scipy.signal](#)

More

We'll use scipy functions based on their need in our examples


Matplotlib

- Plotting library for Python
- Works well with Numpy
- Syntax similar to Matlab

First example

```
import matplotlib.pyplot as plt
import numpy as np
x = np.linspace(0, 10, 1000)
y = np.power(x, 2)
plt.plot(x, y)
```

Result



Modifying the plot


```
plt.xlim((1, 5)) # range for the x axis
plt.ylim((0, 30)) # range for the y axis
plt.xlabel('my x label') # x label
plt.ylabel('my y label') # y label
plt.title('This is the plot title') # plot title
plt.savefig('line_plot_plus2.pdf') #saving the plot in a file
```

Scatterplot

- Use the third parameter to change the format, see details in https://matplotlib.org/stable/api/_as_gen/matplotlib.pyplot.plot.html
- For example “**bo**” indicates blue circles, “**go**” green circles

Example

```
y1 = np.power(x, 2)  
y2 = np.power(x, 3)  
plt.plot(x, y1, "go")  
plt.plot(x, y2, "bo")
```



Histogram

```
import numpy as np
import matplotlib.pyplot as plt

data = np.random.randn(1000) # 1000 random numbers

# creates a figure with subplots, 1 row, 2 columns
# the result is composed of 3 objects, the figure (f1) and the two subfigures (ax1 and ax2)
f, (ax1, ax2) = plt.subplots(1, 2, figsize=(6,3))


# histogram on the first subfigure, with 30 bins, values normalized
ax1.hist(data, bins=30, normed=True, color='b')

# cumulative histogram
ax2.hist(data, bins=30, normed=True, color='r', cumulative=True)

# title for the whole figure
f.suptitle("my histograms")
```


Result

my histograms



Pandas

- General purpose data manipulation lib for Python
- Makes Python very similar to languages like R or Matlab
- We will mainly see how to handle dataframes

What is a dataframe?

- A dataset consisting of rows and columns
- Columns are heterogeneous
- Each column has a name, indicated in the first row
- Rows may have a name, indicated in the first column
- Typically, in machine learning each column of a dataframe is a variable, each row a sample

First steps in Pandas

Creating a data series, similar to an list in Python or in numpy

```
import numpy as np
import pandas as pd
```

```
s = pd.Series([1, 3, 5, np.nan, 6, 8])
```

```
>>>s
```

```
0    1.0
1    3.0
2    5.0
3    NaN
4    6.0
5    8.0
```

```
dtype: float64
```

Creating a Dataframe

Composed of 4 columns, named “A”, “B”, “C”, “D”, each containing 6 random numbers

```
df = pd.DataFrame(np.random.randn(6, 4), columns=list("ABCD"))
```

df

	A	B	C	D
0	0.284643	0.287625	-0.058835	0.422457
1	0.271413	-1.257482	2.048370	0.520126
2	-0.526431	-1.489253	2.163840	-0.329746
3	0.601180	-0.816019	-2.189319	0.713336
4	0.183926	-0.528493	1.498898	-0.287538
5	2.412859	0.211086	-1.003173	-3.027977

Adding row indexes

As said before, not always necessary

```
#a range of six dates
```

```
dates = pd.date_range("20130101", periods=6)
```

```
#dataframe where row indexes are the dates created before
```

```
df = pd.DataFrame(np.random.randn(6, 4), index=dates, columns=list("ABCD"))  
df
```

	A	B	C	D
2013-01-01	-0.297087	1.857124	0.215378	0.560495
2013-01-02	0.481740	-1.332195	0.264884	-0.753363
2013-01-03	-0.032739	-0.722602	-1.857090	1.109269
2013-01-04	0.176271	-0.308361	-1.125783	0.356313
2013-01-05	0.369300	-0.468248	0.218948	-0.700234
2013-01-06	-2.237125	0.889179	1.378012	-0.669981

Creating a dataframe as a dictionary

```
df2 = pd.DataFrame(  
    {  
        "A": 1.0,  
        "B": pd.Timestamp("20130102"),  
        "C": pd.Series(1, index=list(range(4)), dtype="float32"),  
        "D": np.array([3] * 4, dtype="int32"),  
        "E": pd.Categorical(["test", "train", "test", "train"]),  
        "F": "foo",  
    }  
)
```

Result

df2

	A	B	C	D	E	F
0	1.0	2013-01-02	1.0	3	test	foo
1	1.0	2013-01-02	1.0	3	train	foo
2	1.0	2013-01-02	1.0	3	test	foo
3	1.0	2013-01-02	1.0	3	train	foo

Note: you can see how scalar are repeated to fill the dataframe

Column types

```
df2.dtypes
```

```
A          float64
```

```
B    datetime64[ns]
```

```
C          float32
```

```
D          int32
```

```
E          category
```

```
F          object
```

```
dtype: object
```

Utility functions

```
df.head(n) # first n rows  
df.tail(n) # last n rows  
df.columns # column names  
df.index #row names  
df.to_numpy() #convert to numpy array  
df.T #transpose
```

Statistics

```
df.describe()
```

	A	B	C	D
count	6.000000	6.000000	6.000000	6.000000
mean	-0.256606	-0.014184	-0.150942	-0.016250
std	1.009803	1.170434	1.152639	0.797090
min	-2.237125	-1.332195	-1.857090	-0.753363
25%	-0.231000	-0.659013	-0.790493	-0.692671
50%	0.071766	-0.388304	0.217163	-0.156834
75%	0.321043	0.589794	0.253400	0.509450
max	0.481740	1.857124	1.378012	1.109269

Sorting

On the second (horizontal) axis (axis=1), descending

```
df.sort_index(axis=1, ascending=False)
```

	D	C	B	A
2013-01-01	0.560495	0.215378	1.857124	-0.297087
2013-01-02	-0.753363	0.264884	-1.332195	0.481740
2013-01-03	1.109269	-1.857090	-0.722602	-0.032739
2013-01-04	0.356313	-1.125783	-0.308361	0.176271
2013-01-05	-0.700234	0.218948	-0.468248	0.369300
2013-01-06	-0.669981	1.378012	0.889179	-2.237125

Sorting by column

```
df.sort_values(by="B")
```

	A	B	C	D
2013-01-02	0.481740	-1.332195	0.264884	-0.753363
2013-01-03	-0.032739	-0.722602	-1.857090	1.109269
2013-01-05	0.369300	-0.468248	0.218948	-0.700234
2013-01-04	0.176271	-0.308361	-1.125783	0.356313
2013-01-06	-2.237125	0.889179	1.378012	-0.669981
2013-01-01	-0.297087	1.857124	0.215378	0.560495

Selection

```
df["A"] #column "A"  
df[0:2] #first two rows  
df["20130102":"20130104"] #rows with a given data range  
df.loc["20130102":"20130104", ["A", "B"]] #multiple selection on rows and columns  
df.loc[:, ["A", "B"]] #multiple selection on columns only
```

Accessing elements

```
df.loc[dates[0], "A"] #by keys  
df.iloc[1,3] #by index  
df.iloc[1,1:2] #multiple elements
```

Conditional selection

```
df[df["A"] > 0] #on a single column
```

```
df[df > 0] #on all data
```


Adding a new column

```
s1 = pd.Series([1, 2, 3, 4, 5, 6], index=pd.date_range("20130102", periods=6))  
df["F"] = s1
```

Missing data

Pandas uses `np.nan` to represent missing data.

```
#Adds an empty column with missing data
```

```
df1 = df.reindex(index=dates[0:4], columns=list(df.columns) + ["E"])
```

```
#Fills the first 2 lines
```

```
df1.loc[dates[0] : dates[1], "E"] = 1
```

```
df1
```

	A	B	C	D	E
2013-01-01	-0.297087	1.857124	0.215378	0.560495	1.0
2013-01-02	0.481740	-1.332195	0.264884	-0.753363	1.0
2013-01-03	-0.032739	-0.722602	-1.857090	1.109269	NaN
2013-01-04	0.176271	-0.308361	-1.125783	0.356313	NaN

Missing data functions

```
df1.dropna(how="any")  
#removes lines with missing numbers
```

```
df1.fillna(value=5)  
# replaces missing by 5
```

```
pd.isna(df1)  
# returns True or False for any value,  
# indicating whether it is a missing data or not
```

Converting types

For example, from string into categorical

```
df["E"] = pd.Series(["Red", "Green", "Red", "Blue", "Blue", "Green"])
df["E"] = df["E"].astype("category")
```

```
df.dtypes
A      float64
B      float64
C      float64
D      float64
E      category
```

Concatenating data frames

```
df2=pd.concat([df,df])
```

```
#in this cases, the results contains twice the rows of df
```

Stats

```
#applies a function over columns - computes the mean for each column  
df2.apply(np.mean)
```

```
#applies a function over columns - normalizes the values  
df4=df2.apply(lambda x: (x-x.min()) / (x.max() - x.min()))
```

Reading/writing csv file

```
df.to_csv("foo.csv") #writes dataframe to a CSV file
```

```
dfx=pd.read_csv("foo.csv") #reads it
```

More...

<https://pandas.pydata.org>