

# Probabilistic Learning

Naive Bayes Classifiers

# Probability-based predictions

- Weather forecast says things like “70% chance of rain”
- They express probability of precipitation
- How are these probability calculated?

# Use of past data

- Probability-based prediction use past data to produce predictions for future events
- E.g. in the past, under similar conditions, in 7 out of 10 cases it rained

# Naive Bayes algorithm

- Originates from the work of the mathematician Thomas Bayes (18th century)
- Classifiers based on Bayesian methods use training data to compute the probability of each outcome based on evidence from feature values

# Applications

- Text classification
- Intrusion/anomaly detection in computer networks
- Diagnosis of medical conditions from symptoms

# Where should you apply it?

- Many features
- Even if some features alone have little effect, their combination can have a quite large effect on the model

# Event, trials

Bayesian probability theory estimates the likelihood of an **event** based on the evidence of multiple **trials**

Event	Trial
Head result	Coin flip
Rainy weather	A single day's weather
Message is spam	Incoming email message
Candidate becomes president	Presidential election
Win the lottery	Lottery ticket

# From trials to probability


- Probability is estimated by dividing the number of trials in which the event occurred by the total number of trials
- If 10 out of 50 similar messages contained spam, the the probability of spam =  $10/50 = 0.20$  or 20%
- If it rained 3 out of 10 days under similar conditions, the raining probability is  $3/10=0.30$  or 30%



# Complementing probability

- The cumulative probability of all outcomes is 1
- if  $P(\text{spam})=0.20$ , then  $P(\text{ham})=1-0.20=0.80$
- Spam and Ham are **mutually exclusive** events (i.e., they cannot occur at the same time)
- Spam and Ham are **exhaustive**, i.e., together they represent all possible outcomes
- Notation for complement:  $A'$  or  $A^c$
- Probability notation:  $P(A')$  or  $P(A^c)$

# Spam and its complement

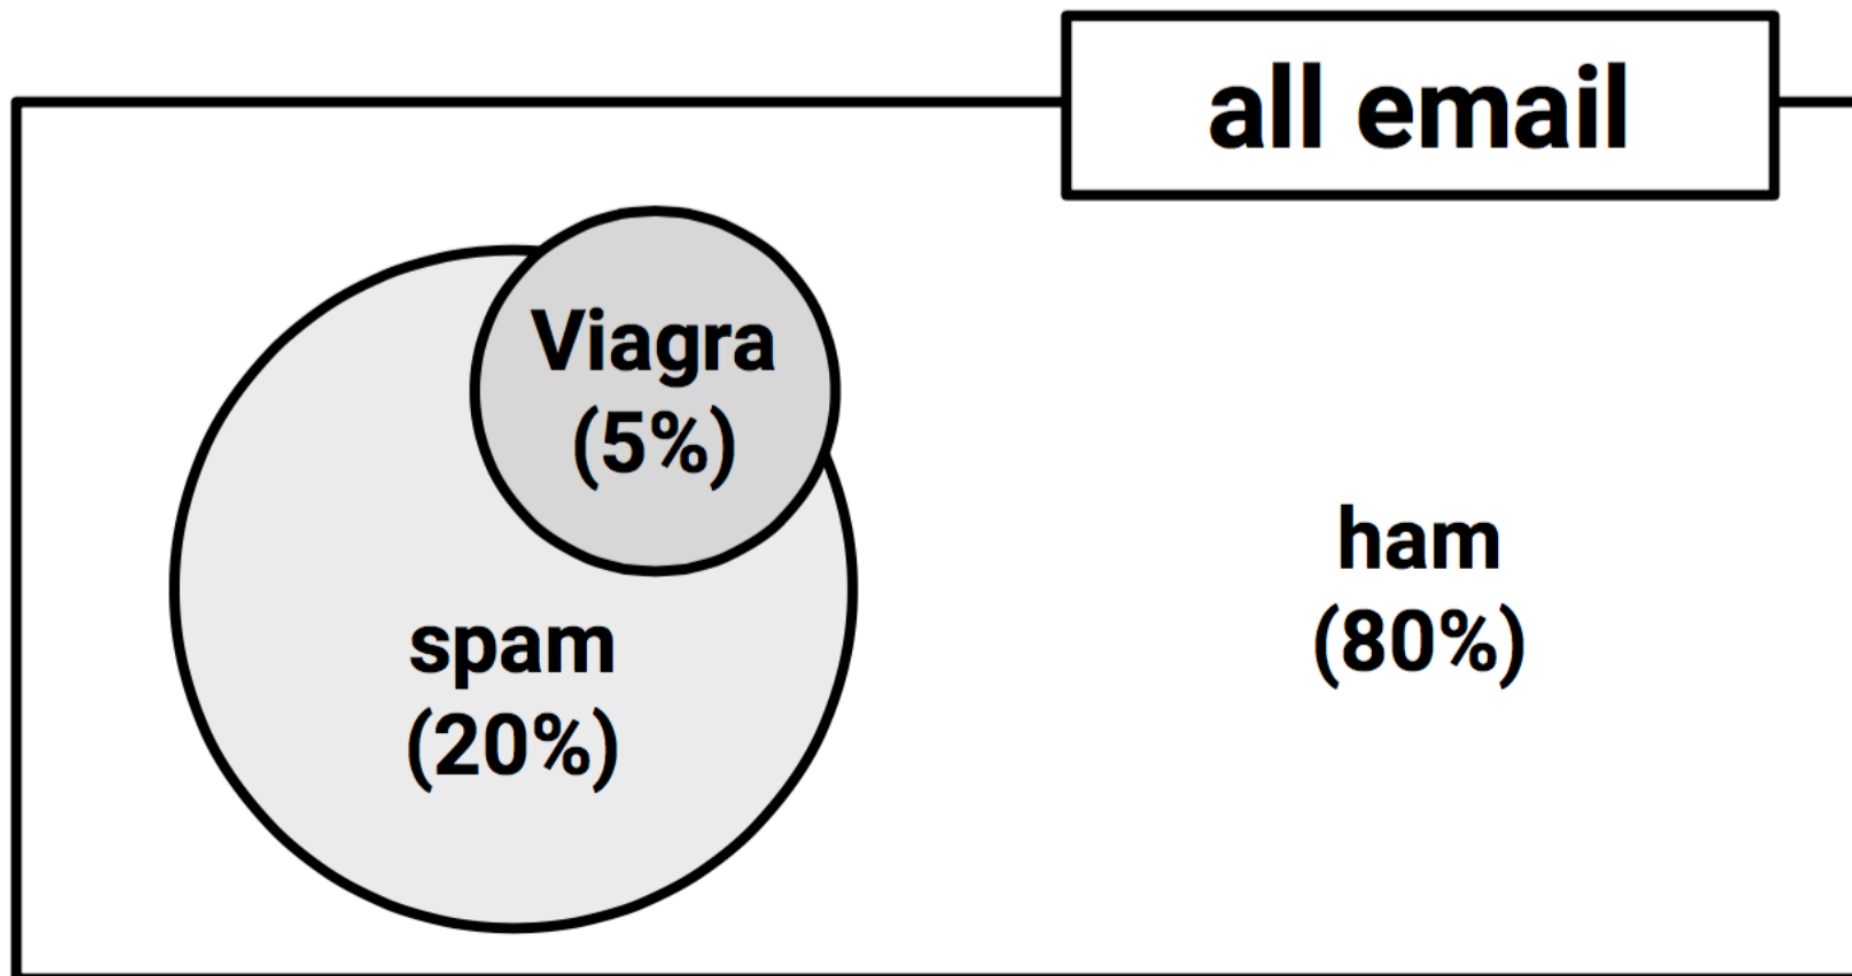


# Joint probability

- Often we are interested to monitor several non-mutually exclusive events for the same trial
- If they occur concurrently, we may exploit them to make predictions

# Example


- Let's consider all email messages that contain the word "Viagra"
- Are they spam?



# Example - discussion

- Not all spam contain the word “Viagra”
- Not all messages with the word “Viagra” are spam
- However, the diagram shows that in “most of the cases” the word “Viagra” appears for spam
- Can we quantify that?

# Venn diagram



# Facts

- 20% of all messages are spam
- 5% of all message contain the word “Viagra”
- What’s the overlap?

# Joint probability

- Probability that both spam and Viagra occur
- Written as  $P(\text{spam} \cap \text{Viagra})$
- Calculating the joint probability depends on whether events are **independent** or **related**




# Independent events

- **Independence** means that knowing the outcome of one event does not provide any information about the other
- For example:
  - the outcome of a coin flip is independent from whether it is rainy or sunny on any given day
  - the outcome of a coin flip is independent from previous coin flips
- If all events are independent, it would be impossible to predict one event by observing another
- Therefore, we exploit **dependent** events for predictive modeling

# Dependent events

- The presence of cloud is predictive of a rainy day
- The word Viagra may be predictive of a spam email



# Probability of dependent events

- If  $P(\text{spam})$  and  $P(\text{Viagra})$  were independent,  $P(\text{spam} \cap \text{Viagra})$  would have been the probability of both events happening at the same time, i.e.

$$P(A \cap B) = P(A) * P(B)$$

- In our case:

$$P(\text{spam} \cap \text{Viagra}) = P(\text{spam}) * P(\text{Viagra}) = 0.20 * 0.05 = 0.01$$

# However...

In our case, the events are likely to be dependent,  
therefore this calculation is incorrect

# The Bayes' theorem

- Estimates the probability of an event conditioned by the evidence provided by another:

$$P(A | B) = \frac{P(A \cap B)}{P(B)}$$

- $P(A|B)$  is said **conditional probability**

# Bayes' Theorem

- $P(A \cap B) = P(A | B) * P(B)$
- also, since  $P(A \cap B) = P(B \cap A)$ :

$$P(A | B) = \frac{P(A \cap B)}{P(B)} = \frac{P(B | A)P(A)}{P(B)}$$

# Prior probability, likelihood, marginal likelihood

- $P(\text{spam})$  is based on the evidence that previous messages are spam, and is known as **prior probability**
  - This is 20% in our case
- The probability that the word “Viagra” was used in previous spam messages,  $P(\text{Viagra}|\text{spam})$  is called the **likelihood**
- The probability that “Viagra” appeared in any message at all,  $P(\text{Viagra})$  is called **marginal likelihood**

# Probability vs. Likelihood

- The likelihood means increasing the chances of a particular situation to occur by varying a certain distribution (i.e., how would the chances of a message to be spam increase if the word “Viagra” is there)?
- Probability simply gives you a view of the current distribution, indicating the chances of something to happen.



# Posterior probability

- We can compute a posterior probability that a message is spam given that it contains the word Viagra:

The diagram illustrates the calculation of the posterior probability  $P(\text{spam} \mid \text{Viagra})$  using Bayes' theorem. The equation is presented as 
$$P(\text{spam} \mid \text{Viagra}) = \frac{P(\text{Viagra} \mid \text{spam})P(\text{spam})}{P(\text{Viagra})}$$
. Four arrows provide context: an arrow from 'Likelihood' points to  $P(\text{Viagra} \mid \text{spam})$ ; an arrow from 'Prior Probability' points to  $P(\text{spam})$ ; an arrow from 'Marginal Likelihood' points to  $P(\text{Viagra})$ ; and an arrow from 'Posterior Probability' points to  $P(\text{spam} \mid \text{Viagra})$ .

$$P(\text{spam} \mid \text{Viagra}) = \frac{P(\text{Viagra} \mid \text{spam})P(\text{spam})}{P(\text{Viagra})}$$

Labels and arrows:

- Likelihood** points to  $P(\text{Viagra} \mid \text{spam})$
- Prior Probability** points to  $P(\text{spam})$
- Marginal Likelihood** points to  $P(\text{Viagra})$
- Posterior Probability** points to  $P(\text{spam} \mid \text{Viagra})$

# Calculating the components of the Bayes theorem

We first construct a frequency table

	Viagra		
Frequency	Yes	No	Total
spam	4	16	20
ham	1	79	80
Total	5	95	100

# Calculating the components of the Bayes theorem

Then, we derive a likelihood table from it, computing the conditional probabilities

	Viagra		
Frequency	Yes	No	Total
spam	4	16	20
ham	1	79	80
Total	5	95	100

# Likelihood table

	Viagra		
Frequency	Yes	No	Total
spam	4	16	20
ham	1	79	80
Total	5	95	100

	Viagra		
Likelihood	Yes	No	Total
spam	4 / 20	16 / 20	20
ham	1 / 80	79 / 80	80
Total	5 / 100	95 / 100	100

# Likelihood table: discussion

Likelihood	Viagra		Total
	Yes	No	
spam	4 / 20	16 / 20	20
ham	1 / 80	79 / 80	80
Total	5 / 100	95 / 100	100

- $P(\text{Viagra} = \text{Yes} | \text{spam}) = 4/20 = 0.20$ , i.e. 20% probability that a message contains “Viagra” when it is spam
- $P(\text{spam} \cap \text{Viagra}) = P(\text{Viagra} | \text{spam}) * P(\text{spam}) = (4/20) * (20/100) = 0.04$

This means that 4 out of 100 spam messages were spam with the term viagra

This is 4 times higher than 0.01, the previous combined probability we computed under the false assumption of independence

# Computing the posterior probability

$$P(\text{spam}|\text{Viagra}) = P(\text{Viagra}|\text{spam}) * P(\text{spam}) / P(\text{Viagra}) = (4/20) * (20/100) / (5/100) = 0.80$$

- Therefore there is 80% probability that a message is spam given that it contains the word “Viagra”
- Therefore, we can use such a word as a reliable predictor
- This is more or less like modern spam filters work...

# The Naive Bayes Algorithm

# Main assumption

- it based on the “naive” assumption that **all features in the data are equally important and independent**
- This is rarely true in real-world
- For example, the email sender may be a more important indicator of spam than any word in the message text



# Strengths and Weaknesses

## Strengths

Simple, fast, yet effective

Works well with noisy and missing data

Works well with a small training, but also with large ones

Estimated probability for a prediction easy to obtain

## Weaknesses

The underlying assumption is often faulty

Not ideal if you have many numeric features

Estimated probabilities less reliable than predicted classes

# What if we have multiple features for prediction?

- Before we only considered the term “Viagra”
- What if we also have other terms, such as “money”, “groceries”, “unsubscribe”
- Let’s compute a new likelihood table...

# Likelihood table...

Likelihood	Viagra ( $W_1$ )		Money ( $W_2$ )		Groceries ( $W_3$ )		Unsubscribe ( $W_4$ )		Total
	Yes	No	Yes	No	Yes	No	Yes	No	
spam	4 / 20	16 / 20	10 / 20	10 / 20	0 / 20	20 / 20	12 / 20	8 / 20	20
ham	1 / 80	79 / 80	14 / 80	66 / 80	8 / 80	71 / 80	23 / 80	57 / 80	80
Total	5 / 100	95 / 100	24 / 100	76 / 100	8 / 100	91 / 100	35 / 100	65 / 100	100

# Computing the posterior probability

- Let's call the words  $W_1, W_2, W_3, W_4$
- For example, let's compute the probability that a message is spam given that it contains: Viagra=Yes, Money=No, Groceries=No, Unsubscribe=Yes

$$P(\text{spam} | W_1 \cap W'_2 \cap W'_3 \cap W_4) = \frac{P(W_1 \cap W'_2 \cap W'_3 \cap W_4 | \text{spam}) P(\text{spam})}{P(W_1 \cap W'_2 \cap W'_3 \cap W_4)}$$

- This formula is computationally intensive to solve, you can imagine what happens if we add more words...

# Moreover...

If an event (a word) was not observed in past data, it will result in a zero probability, and the product will become zero!

# Using Naive assumption

- All predicting features are independent from each other
- Therefore, the numerator of our formula becomes a simple multiplication of probabilities rather than a complex, conditional joint probability
- Finally, the denominator does not depend on the target class (spam or ham), therefore it can be treated as a constant value and hence ignored

# Therefore

- The conditional probability of spam is:

$$P(spam | W_1 \cap W'_2 \cap W'_3 \cap W_4) \propto P(W_1 | spam) P(W'_2 | spam) P(W'_3 | spam) P(W_4 | spam) P(spam)$$

- Similarly, the conditional probability of ham is:

$$P(ham | W_1 \cap W'_2 \cap W'_3 \cap W_4) \propto P(W_1 | ham) P(W'_2 | ham) P(W'_3 | ham) P(W_4 | ham) P(ham)$$

**We ignore the denominator as it is independent on spam or ham**

# Applying it

	Viagra ( $W_1$ )		Money ( $W_2$ )		Groceries ( $W_3$ )		Unsubscribe ( $W_4$ )		
Likelihood	Yes	No	Yes	No	Yes	No	Yes	No	Total
spam	4 / 20	16 / 20	10 / 20	10 / 20	0 / 20	20 / 20	12 / 20	8 / 20	20
ham	1 / 80	79 / 80	14 / 80	66 / 80	8 / 80	71 / 80	23 / 80	57 / 80	80
Total	5 / 100	95 / 100	24 / 100	76 / 100	8 / 100	91 / 100	35 / 100	65 / 100	100

Viagra=Yes, Money=No, Groceries=No, Unsubscribe=Yes

- Overall likelihood of spam:  
 $(4/20) * (10/20) * (20/20) * (12/20) * (20/100) = 0.012$
- Overall likelihood of ham:  
 $(1/80) * (66/80) * (71/80) * (23/80) * (80/100) = 0.002$
- $0.012/0.002=6$ , hence this message is six times more likely to be spam than ham



# Converting likelihoods to probabilities

- We must reintroduce the denominator that we ignored before
- The denominator is the sum of all possible likelihoods, i.e., likelihood of ham + likelihood of spam:
  - $P(\text{spam}) = 0.012 / (0.012 + 0.002) = \mathbf{0.857}$
  - $P(\text{ham}) = 0.002 / (0.012 + 0.002) = 0.143$

# On summary

Given the found pattern, we expect the message to be spam with 85.7% probability

# Generalized formula

The probability of level L for class C, given the evidence of features  $F_1, F_n$ , is equal to the product of

- probabilities of each piece of evidence conditioned on the class level
- the prior probability of the class level  $p(C_L)$
- a scaling factor  $1/Z$  which converts likelihood into probabilities

$$P(C_L | F_1, \dots, F_n) = \frac{1}{Z} p(C_L) \prod_{i=1}^n p(F_i | C_L)$$

# Laplace Estimator

Let's assume we have a new message containing all four terms: "Viagra", "groceries", "money", and "Unsubscribe"

# Let's compute $p(\text{spam})$ and $p(\text{ham})$

Likelihood	Viagra ( $W_1$ )		Money ( $W_2$ )		Groceries ( $W_3$ )		Unsubscribe ( $W_4$ )		Total
	Yes	No	Yes	No	Yes	No	Yes	No	
spam	4 / 20	16 / 20	10 / 20	10 / 20	0 / 20	20 / 20	12 / 20	8 / 20	20
ham	1 / 80	79 / 80	14 / 80	66 / 80	8 / 80	71 / 80	23 / 80	57 / 80	80
Total	5 / 100	95 / 100	24 / 100	76 / 100	8 / 100	91 / 100	35 / 100	65 / 100	100

- Using the likelihood table, we compute the likelihood of spam:  
 $(4/20) * (10/20) * (0/20) * (12/20) * (20/100) = 0$

- Similarly, the likelihood of ham:  
 $(1/80) * (14/80) * (8/80) * (23/80) * (80/100) = 0.00005$

- The probability of spam is:  
 $0 / (0 + 0.00005) = 0$

- The probability of ham is:  
 $0.00005 / (0 + 0.00005) = 1$

# This doesn't make sense!

- This message won't be classified as spam, but its content may suggest it is spam, instead
- This problem arises if an event never occurred in our training set
- In our case the term “groceries” never appeared in spam messages, therefore  $P(\text{spam}|\text{groceries})=0\%$
- Because probabilities in the Naive Bayes formula are multiplied, a zero for a feature results in a zero probability!

# Solution:

## the Laplace estimator

- Idea: add a small number to each frequency, so the multiplication would never become zero
- In practice, this can be any number
- For a large training set a Laplace value=1 is often used

# Let's recompute the number

Likelihood	Viagra ( $W_1$ )		Money ( $W_2$ )		Groceries ( $W_3$ )		Unsubscribe ( $W_4$ )		Total
	Yes	No	Yes	No	Yes	No	Yes	No	
spam	4 / 20	16 / 20	10 / 20	10 / 20	0 / 20	20 / 20	12 / 20	8 / 20	20
ham	1 / 80	79 / 80	14 / 80	66 / 80	8 / 80	71 / 80	23 / 80	57 / 80	80
Total	5 / 100	95 / 100	24 / 100	76 / 100	8 / 100	91 / 100	35 / 100	65 / 100	100

Since we have 4 features (4 words) we add 1 to all numerators of our formula, but we have to add 4 to the denominators



# Let's recompute the number

Likelihood	Viagra ( $W_1$ )		Money ( $W_2$ )		Groceries ( $W_3$ )		Unsubscribe ( $W_4$ )		Total
	Yes	No	Yes	No	Yes	No	Yes	No	
spam	4 / 20	16 / 20	10 / 20	10 / 20	0 / 20	20 / 20	12 / 20	8 / 20	20
ham	1 / 80	79 / 80	14 / 80	66 / 80	8 / 80	71 / 80	23 / 80	57 / 80	80
Total	5 / 100	95 / 100	24 / 100	76 / 100	8 / 100	91 / 100	35 / 100	65 / 100	100

- Likelihood of spam:  
 $(5/24) * (11/24) * (1/24) * (13/24) * (20/100) = 0.0004$
- Similarly, the likelihood of ham:  
 $(2/84) * (15/84) * (9/85) * (24/84) * (80/100) = 0.0001$
- The probability of spam is:  
 $0.0004 / (0.0004 + 0.0001) = 0.80$
- The probability of ham is:  
 $0.0001 / (0.0004 + 0.0001) = 0.20$

**This makes  
more sense!**

# Note

- We added the Laplace estimator (1) to each numerator and to the denominators of the likelihood, we did not add it to the prior probabilities
- Therefore, they remained equal to 20/100 and 80/100
- This is because they are just computed based on observed data and on such adjustment is necessary

# Using numeric features with Naive Bayes

- With Naive Bayes, we use frequency tables from learning over the data
- Therefore, each feature must be categorical to create the likelihood table


**What to do with numeric features?**

# Solution: discretize

- Similar to what done in histograms, we split numerical data into bins and for each bin we create a different categories
- How to split into bins?
- Not really like histogram, but we can use natural splits (e.g., hours if the variable is time), or points where we observe the variable pattern changes quite a bit

# Example

- Assuming we have the time of the day when emails arrive
- If we plot the histogram of this time:



# Discussion

- We notice how the frequency of email varies between night, morning, afternoon, and evening
- Therefore, for each mail we might have four Boolean variables indicating the time band in which the email has arrived (only one of them will be True each time)

# Example

Classifying spam



# Dataset

- SMS spam collection:  
<https://www.dt.fee.unicamp.br/~tiago/smsspamcollection/>

# Examples

## Ham

Better. Made up for Friday and stuffed myself like a pig yesterday. Now I feel bleh. But at least its not writhing pain kind of bleh.

If he started searching he will get job in few days. he have great potential and talent.

## Spam

Congratulations ur awarded 500 of CD vouchers or 125gift guaranteed & Free entry 2 100 wkly draw txt MUSIC to 87066

December only! Had your mobile 11mths+? You are entitled to update to the latest colour camera mobile for Free! Call The Mobile Update Co FREE on 08002986906

# Reading and examining the dataset

```
import pandas as pd
from nltk.corpus import stopwords
import gensim
import numpy as np

dataset=pd.read_csv("sms_spam.csv")

print(dataset.head())
print ("Shape:", dataset.shape, '\n')
```

# Output

```
      type      text
0    ham  Hope you are having a good week. Just checking in
1    ham                                K..give back my thanks.
2    ham          Am also doing in cbe only. But have to pay.
3  spam  complimentary 4 STAR Ibiza Holiday or £10,000 ...
4  spam  okmail: Dear Dave this is your final notice to...
Shape: (5559, 2)
```

# Function for text preprocessing

- This time we will use the out-of-box features available in gensim

# Preprocessing function

```
def transformText(text):
    stops = set(stopwords.words("english"))
    # Convert text to lowercase
    text = text.lower()
    # Strip multiple whitespaces
    text = gensim.corpora.textcorpus.strip_multiple_whitespaces(text)
    # Removing all the stopwords
    filtered_words = [word for word in text.split() if word not in stops]
    # Preprocessed text after stop words removal
    text = " ".join(filtered_words)
    # Remove the punctuation
    text = gensim.parsing.preprocessing.strip_punctuation(text)
    # Strip all the numerics
    text = gensim.parsing.preprocessing.strip_numeric(text)
    # Removing all the words with < 3 characters
    text = gensim.parsing.preprocessing.strip_short(text, minsize=3)
    # Strip multiple whitespaces
    text = gensim.corpora.textcorpus.strip_multiple_whitespaces(text)
    # Stemming
    return gensim.parsing.preprocessing.stem_text(text)
```

# Preprocessing...

```
#applies transformText to all rows of text  
dataset['text'] = dataset['text'].map(transformText)  
print(dataset['text'].head())
```

# Output

```
0             hope good week check
1             give back thank
2             also cbe onli pai
3  complimentari star ibiza holiday cash need urg...
4  okmail dear dave final notic collect tenerif h...
Name: text, dtype: object
```



# Creating training and test set

- In real life, you train your machine learning model on data for which you know the label (in this case “ham” or “spam”)
- When validating an algorithm, we have a dataset with labels everywhere
- We cannot train and test the algorithm on the same data!

# Creating training and test set

- We split the dataset into 2 portions
- **Training set:** used to train the machine learning algorithm
- **Test set:** used to evaluate the algorithm's performances

# How?

- Simply sequentially, if the data has a temporal meaning (e.g. COVID data over time)
- Randomly
- Using cross-validation (we will see this later)
- Note: for now we will split training and test set once, in practice it needs to be done multiple times, and performances analyzed over multiple runs

# Creating a random split using scikit-learn

```
## Split the data  
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(dataset['text'], dataset['type'],  
                                                    test_size=0.33, random_state=10)  
  
print ("Training Sample Size:", len(X_train), ' ', "Test Sample Size:" ,len(X_test))
```

**Training Sample Size: 3724   Test Sample Size: 1835**

# Notes

- Note: the function takes as input the size of the test set and a random number seed
- Initializing the random number seed with a constant always produces the same results
- Initializing the latter with a timestamp makes your results always different

```
from sklearn.model_selection import train_test_split
import time
X_train, X_test, y_train, y_test = train_test_split(dataset['text'], dataset['type'],
                                                    test_size=0.33, random_state=int(time.time()))
```

# Creating a tf-idf corpus for the training set

```
#Build the counting corpus
from sklearn.feature_extraction.text import CountVectorizer
count_vect = CountVectorizer()
X_train_counts = count_vect.fit_transform(X_train)

## Get the TF-IDF vector representation of the data
from sklearn.feature_extraction.text import TfidfTransformer
tfidf_transformer = TfidfTransformer()
X_train_tfidf = tfidf_transformer.fit_transform(X_train_counts)
print ('Dimension of TF-IDF vector :', X_train_tfidf.shape)
```

**Dimension of TF-IDF vector : (3724, 7007)**

# Creating the classifier

```
#Creating the classifier  
#MultinomialNB accepts weights instead of Boolean  
from sklearn.naive_bayes import MultinomialNB  
clf = MultinomialNB()  
# the fit() function of any classifier takes the features from the  
# training set X_train_tfidf and the labels from the training set  
# y_train  
clf.fit(X_train_tfidf, y_train)
```

# Notes

- The `fit()` function applies to any classifier
- It takes as input:
  - the training set features
  - the labels from the training set



# Notes

- Scikit-learn has different kinds of Naive Bayes classifier
  - **BernoulliNB**: accepts boolean as the examples we have seen previously
  - **MultinomialNB**: models counts
  - **GaussianNB**: useful for decimal values provided that they follow a normal distribution
- When creating the models, parameters with options are available, among others the alpha parameter specifies the value of the Laplace estimator (default=1)
  - `clf = MultinomialNB(alpha=0) # no Laplacian smoothing`
  - `clf = MultinomialNB(alpha=1) # smoothing=1`

# Folding the test set into the vector space...

*#indexing the test set*

```
X_new_counts = count_vect.transform(X_test)
X_new_tfidf = tfidf_transformer.transform(X_new_counts)
```

- The `transform()` function takes a set of words and folds them into a given vector space
- We first fold them into `count_vect` (obtained through the `CountVectorizer`) and then we fold the results into the `tfidf_transformer` vector space

# Making the prediction

```
#performing the actual prediction  
predicted = clf.predict(X_new_tfidf)
```

- the `predict()` function is a standard function of different scikit-learn models
- It has to be applied to previously-created model (`clf`)
- It takes as input the features of the test set, folded into the training set space

# Showing the results

```
print(predicted)  
print(np.mean(predicted==y_test))
```

```
['ham' 'ham' 'ham' ... 'spam' 'ham' 'spam']  
0.9607629427792915
```

# Note

- `predicted==y_test` compares every element of predicted (which can be “Ham” or “Spam”) with every element of the test set labels `y_test` (again, “Ham” or “Spam”) and returns True or False
- `np.mean()` computes the fraction of the True over the total, i.e., like computing a mean of 0 and 1 values
- What we computed is called accuracy
- We will later formally define the performance measurements for information retrieval and machine learning

# Exercises

- Try how different models produce different performance values
- For example:
  - Different values of the Laplace estimator
  - Using a BernoulliNB after having indexed the words using a Boolean model:  
`count_vect = CountVectorizer(binary=True)`
  - Filtering out words appearing rarely (try with different values  
`count_vect = CountVectorizer(min_df=10)`)
  - Just using CountVectorizer instead of TfidfTransformer()

# Coming up next...

- Evaluating the model
- Improving the model, by applying feature selection, or by using alternative machine learners
- Applying to sentiment analysis